

COMS4995W31

Applied Machine Learning

Dr. Spencer W. Luo

Columbia University | Spring 2026



Announcement





Mid-term 😐

Week 6

Duration: 1.5 hours

Format:

- Multiple Choice 
- Short Answer 

Cheat Sheet:

- 1-page A4 (double-sided allowed)
- Handwritten or printed



What You Need to Know



Course 1 - 5:

- Data prep & feature engineering
- Generative vs. Discriminative models
- Model evaluation & bias - variance
- Ensemble

Focus on:

- Understanding concepts
- Explaining scenarios
- Sharing intuitions



Assignment 1

<https://www.gradescope.com/courses/1260107/assignments/7670023>

Due: Feb 25, 2026 11:59 PM



Recitation 1

Next week

More details will be released on Ed soon





Model Evaluation

Bias-Variance



AI of the Week: The Death of Public Benchmarks



By early 2026, classic benchmarks like [MMLU](#) and [GSM8K](#) have become effectively useless

Top models (OpenAI, Claude, Gemini) now score >99% on them, yet often fail in real-world edge cases

The Problem: "[Data Contamination](#)"

Since these benchmarks are public on the internet, models have accidentally "trained" on the test questions

The Solution: The rise of Private Leaderboards (e.g., HuggingFace Private). These use dynamic, held-out datasets that model builders never see, ensuring the score reflects true intelligence, not memorization

 Class Connection:

[Overfitting](#)

When a model scores 98% on MMLU but fails your specific task, it is overfitting to the benchmark

It has memorized the "test set" (public internet data) instead of learning the underlying pattern



Agenda

- Motivation
- Train / Validation / Test Split
- Common Evaluation Metrics
- Bias - Variance Tradeoff
- Model Selection Strategies



Motivation



Supervised vs. Unsupervised Learning

Supervised Learning

- Learn mapping $f: X \rightarrow Y$
- Given data (x_i, y_i)
- Example: Spam Email (label = spam / ham)

Unsupervised Learning

- Discover hidden structures in data
- Only x_i , no labels
- Example: Customer Segmentation



What is Machine Learning (ML)?

- Learn patterns from **data**
- Make **predictions** on new data
- **Generalize** beyond training



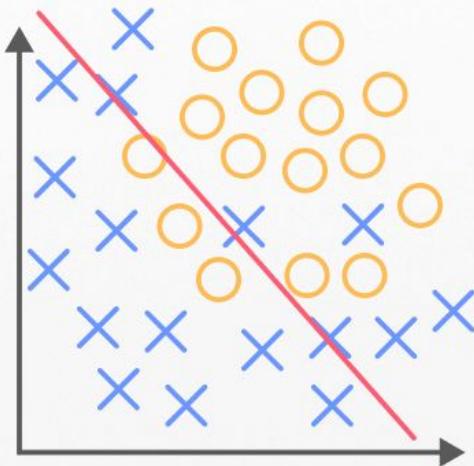
Why Do We Need Model Evaluation? 🤔

Training accuracy ≠ Real-world performance

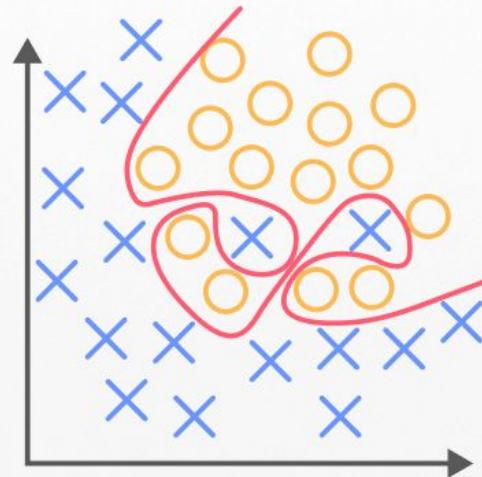
Without proper evaluation:

- Models may memorize training data (**overfitting**)
- Models may be too simple to capture patterns (**underfitting**)

Wrong evaluation → Wrong decisions → Costly mistakes 💰



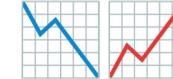
Underfitting



Overfitting



Overfitting vs Underfitting



Overfitting definition:

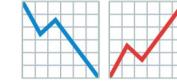
- Model fits the training data too closely - including random noise
- **Symptom:**
 - Training Accuracy $\approx 100\%$
 - Test Accuracy \ll Training Accuracy

Like a student who:

- 📖 Memorizes every word from lecture slides (training set)
- ✗ Cannot answer slightly different questions in mid-term (test set)



Overfitting vs Underfitting



Underfitting definition:

- Model is too simple to capture the underlying patterns in data.
- **Symptom:**
 - Low Training Accuracy
 - Low Test Accuracy

Like a student who:

- ❑ Barely skimmed the textbook
- ✗ Cannot even handle problems in the textbook (training set)
- ✗ Naturally also fails in mid-term (test set)

Why Wrong Metrics Can Be Dangerous



Accuracy is misleading when data is imbalanced

Example: disease detection dataset

- 99% healthy, 1% sick
- A model that always predicts “healthy” → 99% accuracy but useless X

Need metrics that capture what matters: precision, recall, etc.



Key Takeaway



- Evaluation is not optional, it is **essential**
- Good evaluation = Fair model comparison
- Good evaluation = Reliable deployment in real world



Train / Validation / Test Split



Why Split Data? 🍰

Goal: measure **generalization ability**

If we only check training accuracy:

- Might think model is “perfect” → but actually memorizing 📖

Solution: keep separate sets

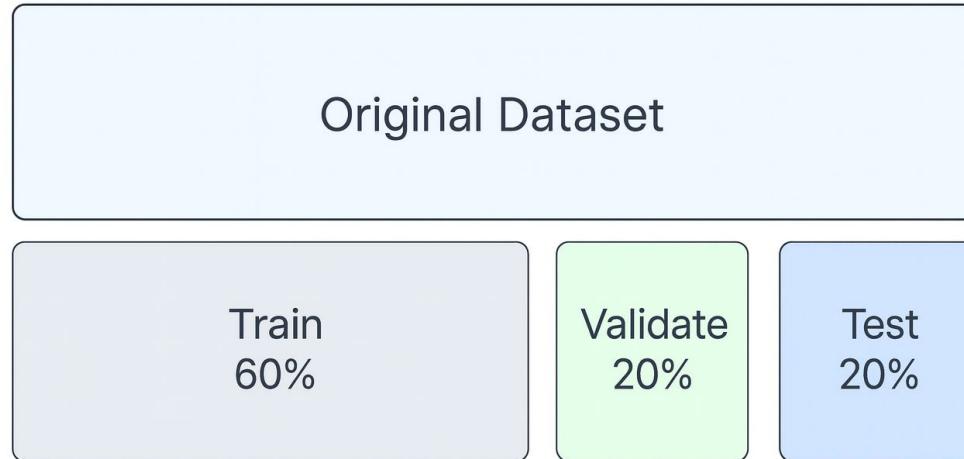
- **Train** → learn patterns
- **Validation** → tune hyperparameters
- **Test** → simulate unseen real-world data



Why Split Data?



Multiple Models Percentage Split





Typical Splits



Common ratios:

- Train: 60 - 80%
- Validation: 10 - 20%
- Test: 10 - 20%

Important rule: Test set is locked





Hold-Out Method



One-time split (Train / Val / Test)

- Pros: simple, fast
- Cons:
 - depends heavily on the random split
 - unstable on small data



k-Fold Cross-Validation



 Idea: Rotate the validation set

Steps:

- Split into k equal folds
- Train on $(k-1)$ folds, validate on the remaining one
- Repeat k times, average results

Pros: robust, uses all data

Cons: more compute

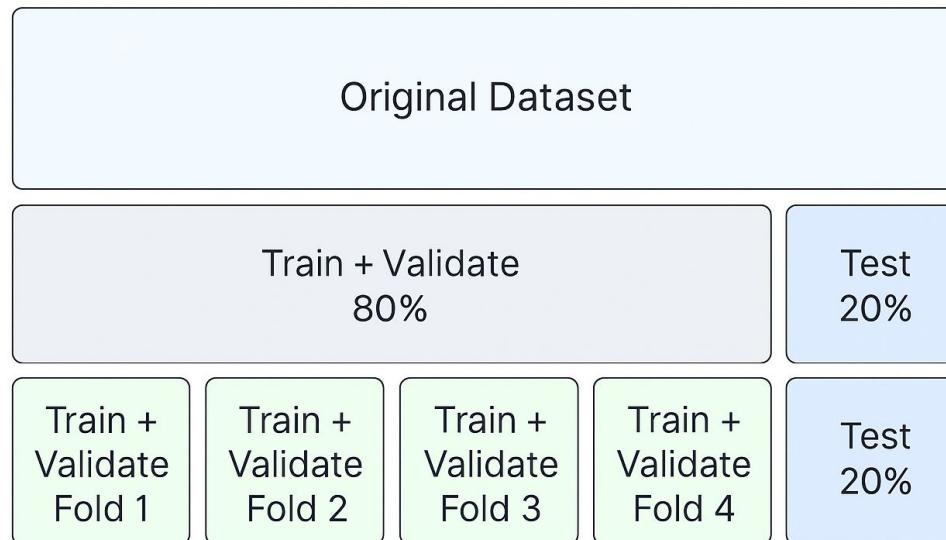




k-Fold Cross-Validation



Multiple Models
K-Fold Cross Validation
 $K = 4$





Common Evaluation Metrics



Why Metrics Matter ⚡

Different problems need different metrics

Example:

- Credit card fraud detection → catch rare cases
- Recommender system → focus on ranking quality

⚠ Accuracy is not enough ❌



Regression Metrics



Mean Squared Error (MSE)

- Penalizes large errors more

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2$$

Mean Error Squared



Regression Metrics



Mean Absolute Error (MAE)

- Less sensitive to outliers

$$MAE = \frac{1}{n} \sum \left| y - \hat{y} \right|$$

Divide by the total number of data points

Predicted output value

Actual output value

Sum of

The absolute value of the residual



Classification Metrics

Accuracy

- Proportion of correct predictions
- Simple but misleading on imbalanced data



Confusion Matrix

1 2
3 4

		Predicted Class			
		Positive	Negative		
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Recall $\frac{TP}{(TP + FN)}$	
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$	
	Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$		



Precision, Recall



Precision (Purity)

- Out of predicted positives, how many are correct?

Recall (Coverage)

- Out of actual positives, how many did we find?

Example: Disease detection



- High recall = fewer missed cases
- High precision = fewer false alarms

F1



F1 combines Precision and Recall into one metric

$$\text{F1 Score} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$



👉 Different Precision & Recall

A model outputs **scores / probabilities**, not just hard labels

To decide class = 1, we must set a threshold

- Example: predict positive if score ≥ 0.5

If we change the threshold:

- Lower threshold \rightarrow more positives predicted \rightarrow higher Recall, lower Precision
- Higher threshold \rightarrow fewer positives predicted \rightarrow higher Precision, lower Recall

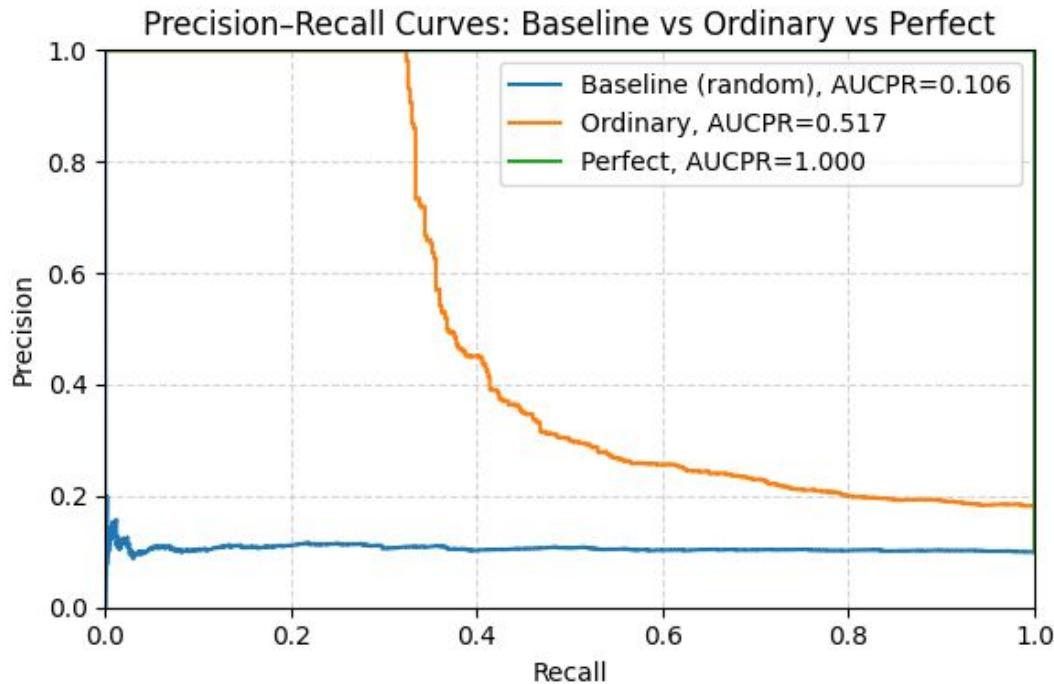
⚖️ This trade-off creates multiple (Precision, Recall) pairs



AUCPR



AUC = Area Under Curve





Choosing the Right Metric



- Regression → MSE, MAE, RMSE
- **Balanced** classification → Accuracy, F1
- **Imbalanced** classification → Precision, Recall, F1, AUCPR
- Unsupervised tasks → Normalized Mutual Information (later lectures)

Metrics



Predicted Class	
Positive	Negative
True Positive (TP)	Sensitivity $\frac{TP}{(TP+FN)}$
Type II Error	Specificity $\frac{TN}{(TN+FP)}$
Precision $\frac{TP}{(TP+FP)}$	False Positive (FP) Type I Error
Precision $\frac{TP}{(TP+FP)}$	Negative Predictive Value $\frac{TN}{(TN+FN)}$



Bias - Variance Tradeoff



Bias–Variance Decomposition



Bias

- Error caused by simplifying assumptions in the model, leading to a systematic difference between prediction and ground truth

Symptom

- Model is too simple to capture the true relationship
- Leads to underfitting
- Example: Using a straight line to fit a curved pattern



Bias–Variance Decomposition

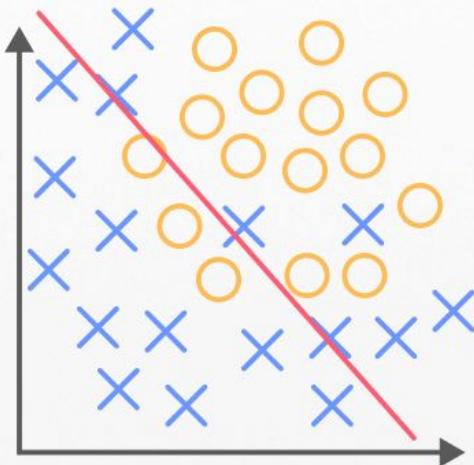


Variance

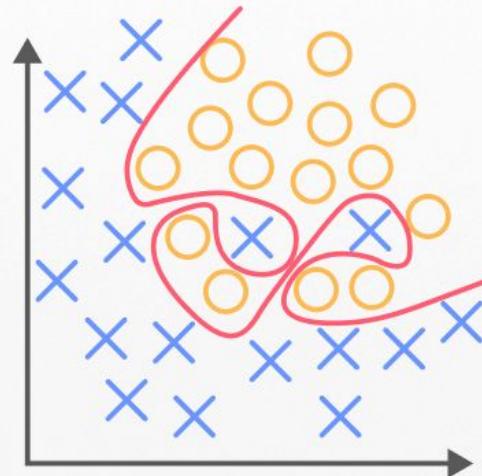
- Error caused by a model's sensitivity to small fluctuations in the training data, leading to inconsistent predictions across different datasets.

Symptom

- Model is too complex, changes a lot with small noises in data
- Leads to overfitting



Underfitting



Overfitting

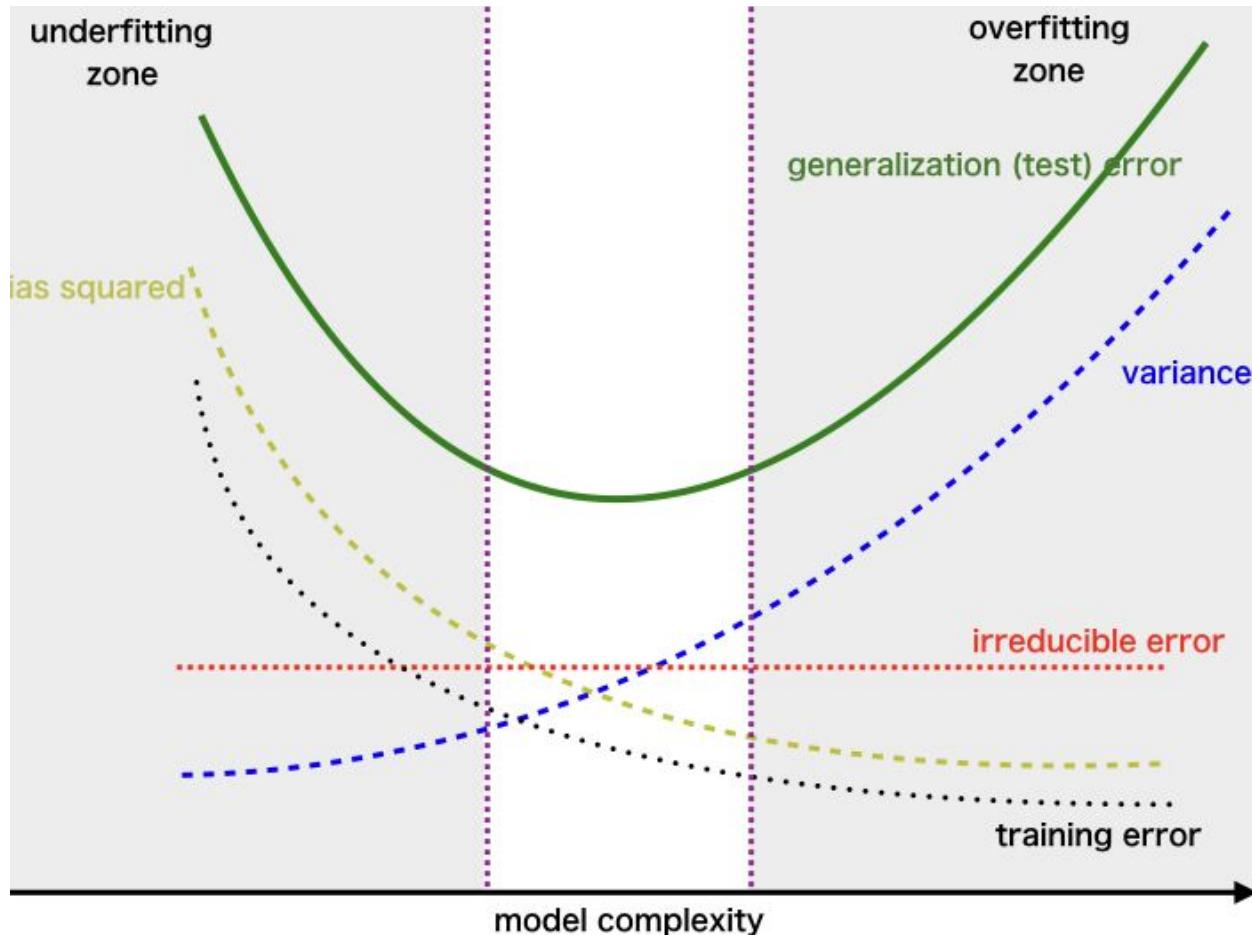


Bias–Variance Tradeoff

Generalization Error = Bias² + Variance + Noise

- Generalization Error == Expected Test Error

Goal = balance between **Bias** (too simple → underfit) and
Variance (too complex → overfit)





Bias vs Variance — Key Takeaways



Opposite failure modes

- High Bias → Underfitting, model too rigid
- High Variance → Overfitting, model too sensitive

Tradeoff is unavoidable

- Reducing bias usually increases variance
- Reducing variance usually increases bias

What we really want

- Minimize expected test error (generalization error)
- Balance both terms + accept some irreducible noise



Model Selection Strategies



Why Model Selection? 🎯

We have mastered:

- How to split data 🍰
- How to measure performance 📊
- Why bias & variance matter ⚖️

Next challenge:

- Choosing **the best model** among candidates



Cross-Validation for Model Choice



Not just for evaluation → also for model selection

Use **k-fold CV** to estimate performance

Pick model with best **average** validation score



Hyper-parameter Tuning



Models often have parameters not learned directly

- Regularization strength (α)
- Decision tree depth
- Neural network depth

Strategies

- Grid search
- Random search





Regularization



Problem: overly complex models → high variance

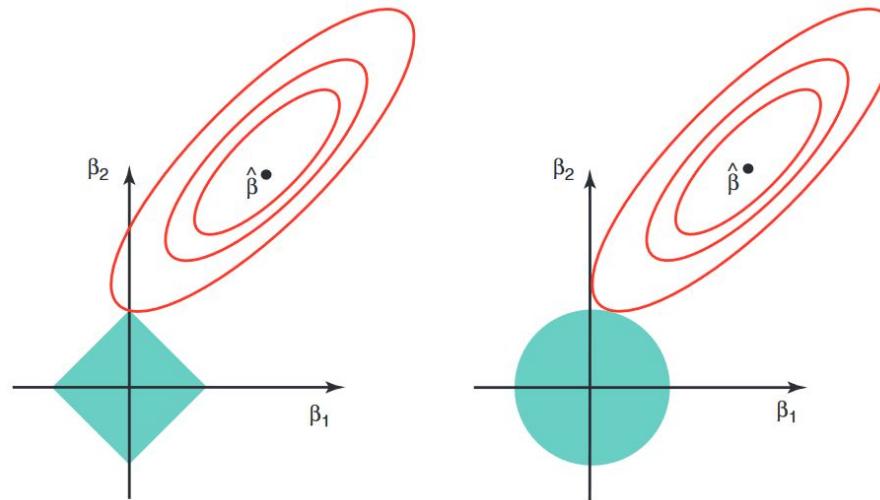
Solution: add penalty terms

- **Ridge (L2)**: penalize large weights smoothly
- **Lasso (L1)**: shrink some weights to zero → feature selection



Ridge vs Lasso (Geometry)

Lasso	Ridge
diamond constraint → corners	circular constraint → smooth shrinkage





Early Stopping



Common in neural networks

Idea:

- Stop training when validation error starts increasing
- Prevents overfitting from too many epochs



Ensemble Methods



Combine multiple models to **reduce variance**

Examples:

- Bagging / Random Forests
- Boosting (AdaBoost, XGBoost)

Analogy: “wisdom of the crowd”



Practical Workflow



- Split data into Train / Val / Test
- Define candidate models + hyperparameters
- Use Cross-validation for fair comparison
- Regularize or stop early if variance too high