

COMS4995W31

Applied Machine Learning

Dr. Spencer W. Luo

Columbia University | Spring 2026





Announcement



Assignment 1

Due Date: Saturday, February 28th, 11:59 PM EST

Ed post: [Link](#)



Midterm

- Monday, March 2, 2026, from 4:10 PM to 5:40 PM
- No cell phones or calculators allowed
- [Sample questions and answers](#)
- Prep: Focus on the lecture slides, raise questions and ask ChatGPT/Gemini for clarifications



Tree Models & Ensembles

AI of the Week: The Undisputed Kings of Tabular Data



Standard structured data (SQL databases) still runs most of the world's real-world business systems

The Reality in 2026:

- Classical Tree Ensembles (XGBoost, LightGBM etc.) remain the absolute gold standard
- [A review of recent Kaggle tabular competitions](#) shows that these gradient boosting frameworks are still the core of almost every single winning solution
- The recent XGBoost 3.1 releases rolled out massive updates
 - introducing blazing-fast GPU-accelerated pipelines and "Native Categorical Re-coding"
 - multimodal support

AI of the Week: The Undisputed Kings of Tabular Data



Ensembles Win

As we will cover today, combining many "weak" decision trees into a powerful ensemble dramatically reduces both bias and variance

Feature Engineering > Raw Compute

Tree models heavily reward clever feature engineering and a deep understanding of data



Agenda

- Motivation
- Decision Trees
- Ensembles
- Summary



Motivation



Why do we need Tree Models? 🌲

In previous:

- **Linear Regression** → assume linear decision boundary
- **Naive Bayes** → assume feature independence

But... real-world data is often non-linear, complex, and mixed



Why do we need Tree Models? 🌲

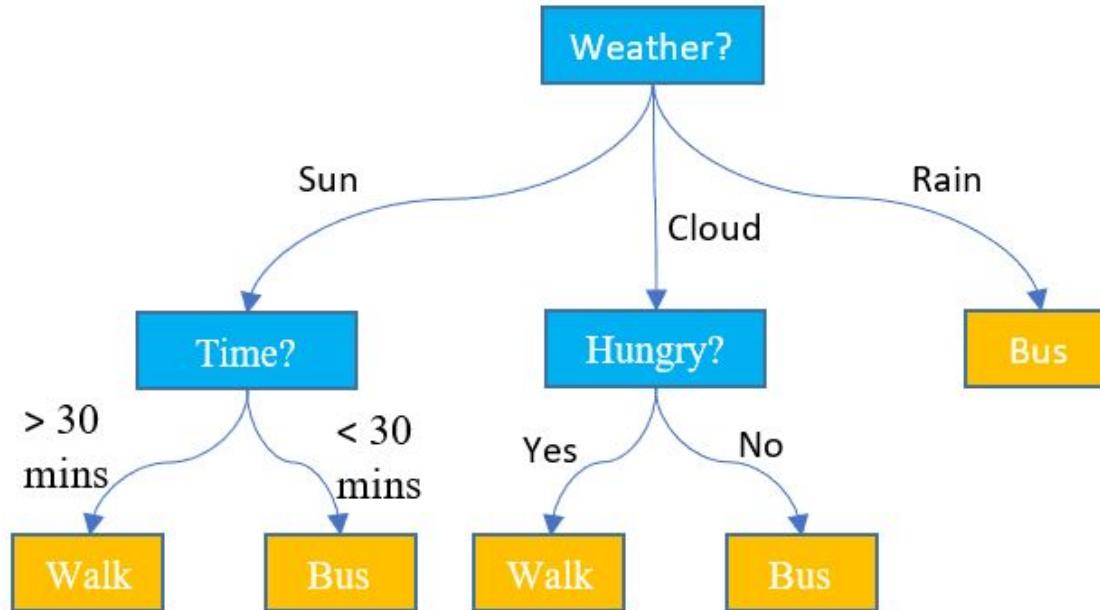
Decision Trees:

- Flexible → no strict distribution or linearity assumptions
- Versatile → handle both **numeric** & **categorical** features
- **Non-linear power** → capture curved/complex decision boundaries
- Interpretability → rules easy to visualize & explain to non-experts
- **Foundation of Ensembles** → Random Forests, Boosting, XGBoost





Why do we need Tree Models? 🌲





Decision Trees



What is a Decision Tree? 🌳

[Inference] A decision tree is like a 20 Questions game 🎲

- At each step → ask a yes/no question about the data
- Each answer leads you further down the tree
- Finally, you reach the final decision / prediction

Example:

Shall we have the mid-term test next week?

- Yes → take mid-term
- No → take mid-term :)



Decision Tree in Action

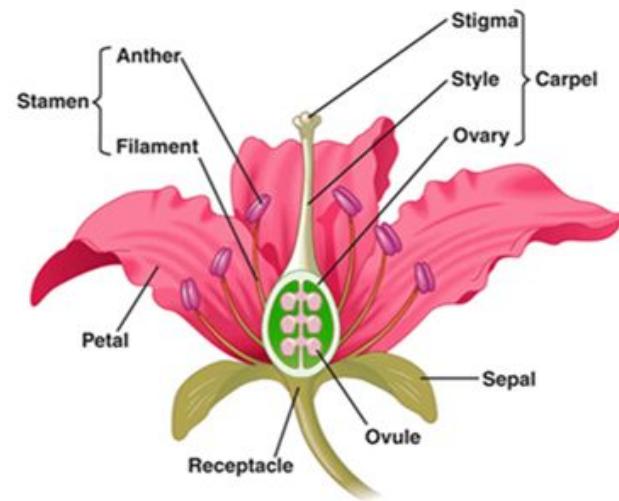
Iris Dataset: 150 flowers → 3 classes

Features:

- Sepal length
- Sepal width
- Petal length
- Petal width

Target labels:

- [0] Iris Setosa
- [1] Iris Versicolor
- [2] Iris Virginica

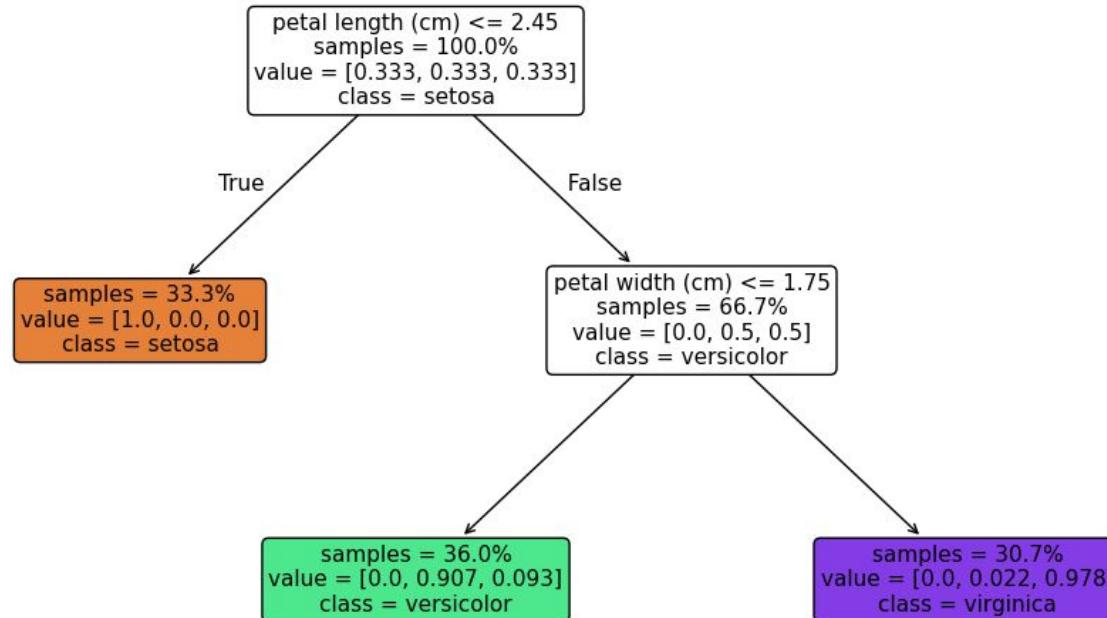




Visualizing a Simple Decision Tree



Iris Decision Tree (max_depth=2)





Components of a Decision Tree



Root Node → initial state

Internal Nodes → decision points (questions based on features)

Branches → outcomes of questions (Yes / No, $>$ / $=$ / $<$)

Leaves → final prediction (class label / value)

Decision Path → sequence of rules from root to leaf

- one classification rule

Tree Depth → longest path from root to leaf (tree complexity)



Two Flavors of Trees



Classification Tree:

- Output = discrete label
- Leaf prediction = **majority class**

Regression Tree:

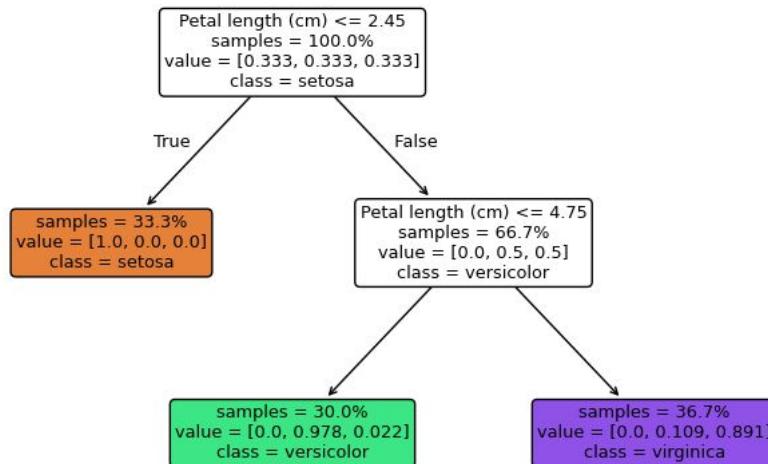
- Output = continuous value
- Leaf prediction = **average of samples**



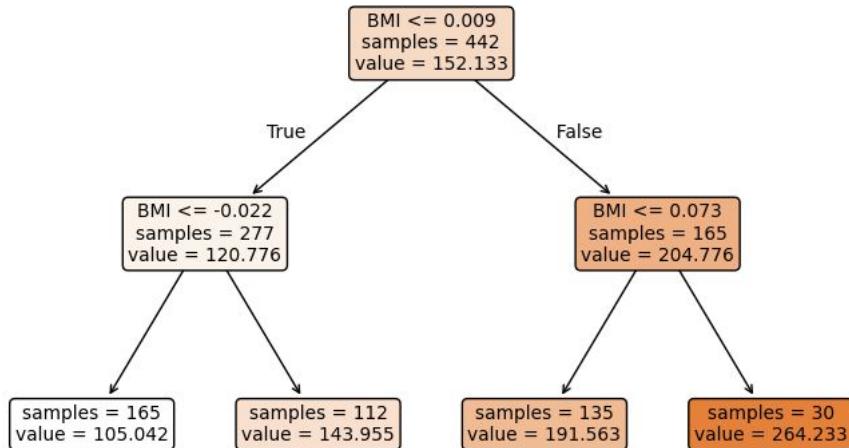
Two Flavors of Trees



Classification Tree
Leaf = Majority Class



Regression Tree
Leaf = Average Value





How do we train (aka split)?

 Idea: Make children nodes “purer” than the parent

Pure node → samples mostly belong to one class

Example intuition:

- Before split:
 - 50% red , 50% blue  (100)
- After split:
 - left node 90% red  (40)
 - right node 80% blue  (60)
- Better separation → better classification capability



Mathematical Support



Pure definition:

Gini Impurity

$$Gini(S) = 1 - \sum_{i=1}^C p_i^2$$

Entropy

$$H(S) = - \sum_{i=1}^C p_i \log p_i$$

p_i : proportion of samples in node S that belong to class i

Smaller \rightarrow Better



Entropy vs Gini

Gini Impurity: How mixed the classes are in a node

Entropy: What the uncertainty degree is in a node

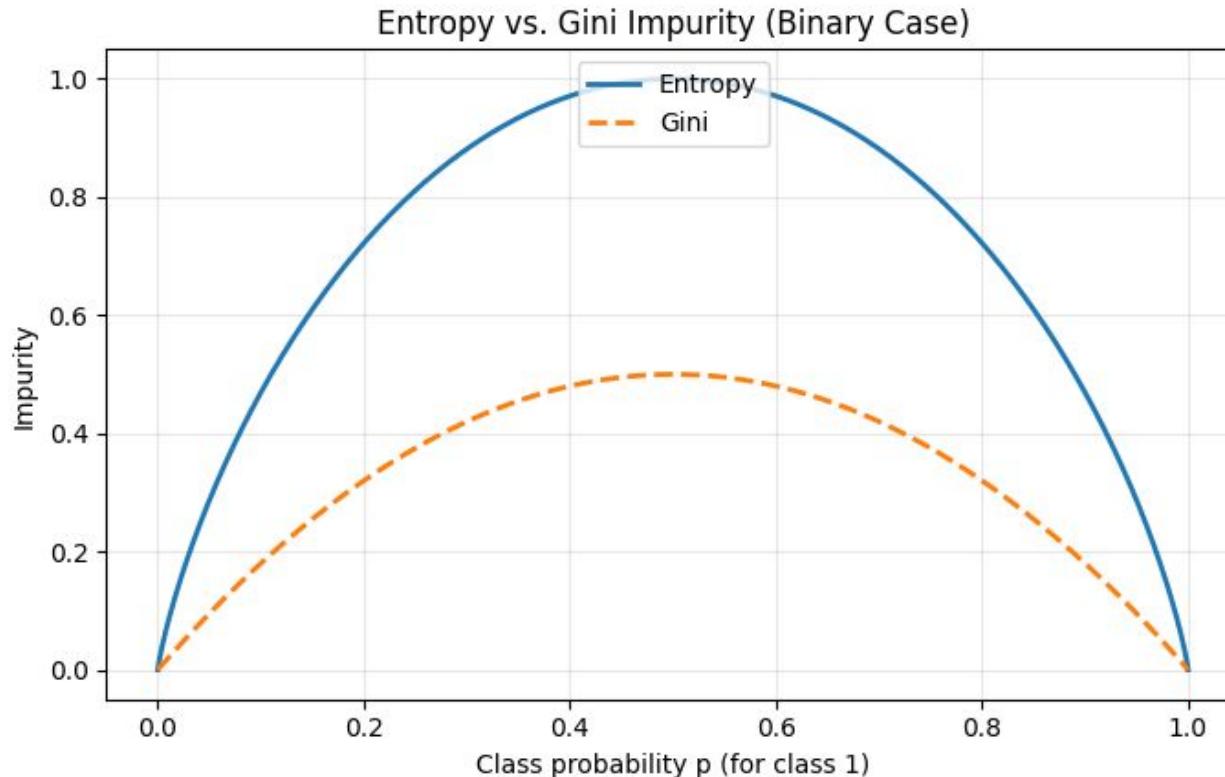
0 = pure, higher = more mixed

Both are similar in shape

Used interchangeably in practice



Entropy vs Gini





Decision Tree Algorithm

```
func build_tree(data) → current node:  
    if stopping_condition(data):  
        return Leaf(prediction=data.label_majority)  
    best_feature, threshold = choose_best_split(data)  
    left_data, right_data = split(data, best_feature, threshold)  
    node = Node(feature=best_feature, threshold=threshold)  
    node.left = build_tree(left_data)  
    node.right = build_tree(right_data)  
    return node
```



How to Split?

```
func choose_best_split(data) → (feature, threshold):  
    best_split = None; best_gain = -inf  
    for feature in features:  
        for threshold in possible_thresholds(feature):  
            left, right = split(data, feature, threshold)  
            if left or right is empty: continue  
            gain = impurity(parent) - weighted_impurity(left, right)  
            if gain > best_gain:  
                best_gain = gain; best_split = (feature, threshold)  
    return best_split
```



When to Stop?

```
func stopping_condition(data, depth):  
    if all_same_label(data):  
        return True  
    if depth >= MAX_DEPTH:  
        return True  
    if split_gain(data) < MIN_GAIN:  
        return True  
# Otherwise, continue splitting  
    return False
```



Deep Trees = Overfitting !

Deep tree memorizes training data → high variance

Small perturbation in data → different splits

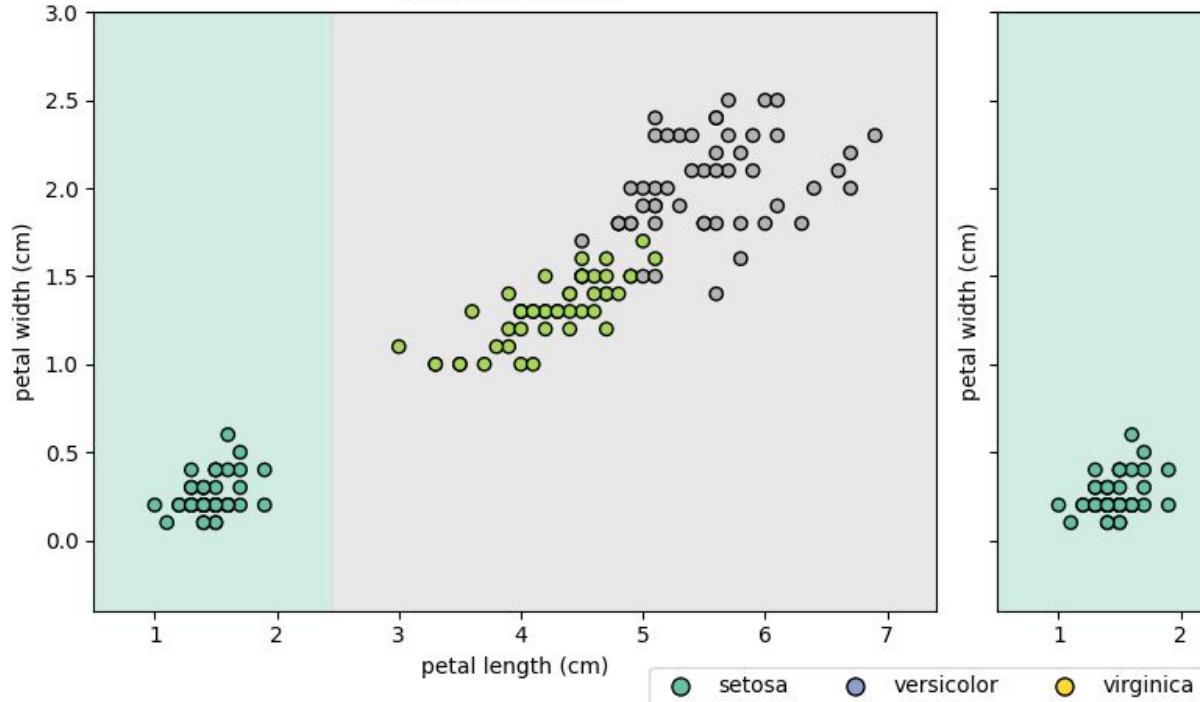
Example: depth=1 vs depth=5

Need control: `max_depth`, `min_samples_split`

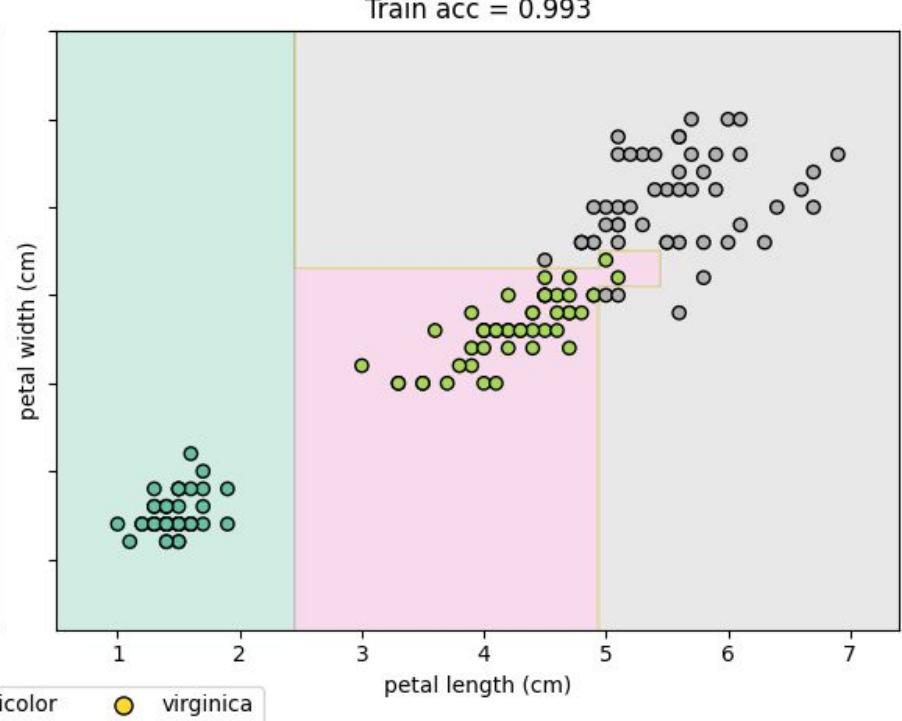


Deep Trees = Overfitting !

Shallow Tree (max_depth=1)
Train acc = 0.667



Deeper Tree (max_depth=5)
Train acc = 0.993





Why Pruning? 🌳剪刀

Without limits, trees grow very deep → low training error, overfitting



Pruning = controlling tree growth to improve generalization

2 perspectives:

- When to prune (Pre vs Post)
- How much to prune (Moderate vs Strong)



Pre-pruning (Early Stopping) When

Put pruning rules **inside** `stopping_condition()`:

- Max depth, min samples, min gain...

Stops growth before overfitting appears

 Fast, simple

 Risk of underfitting if too strict



Post-pruning (Cost-Complexity)



When

Steps:

- First grow a big tree 
- Then cut back unnecessary branches
- Use cross-validation to decide how much to prune

CART (Classification and Regression Trees)

- balance accuracy vs simplicity

 Improves generalization

 Extra computation



Moderate vs Strong Pruning How much

Moderate pruning:

- Keep useful sub-branches
- Balance accuracy & simplicity 😊

Strong pruning:

- Aggressively cut → very shallow tree
- Easier to interpret, but higher bias 😬

Works for both Pre & Post:

- Pre: adjust thresholds (strict vs loose)
- Post: choose small vs large



How to Implement Pruning



- In any decision tree libraries, there is a built-in parameter called `ccp_alpha`
- It controls how much the tree is pruned
- Think of it like **regularization** in linear models:
- Best value is usually chosen by **empirical study** or **cross-validation**

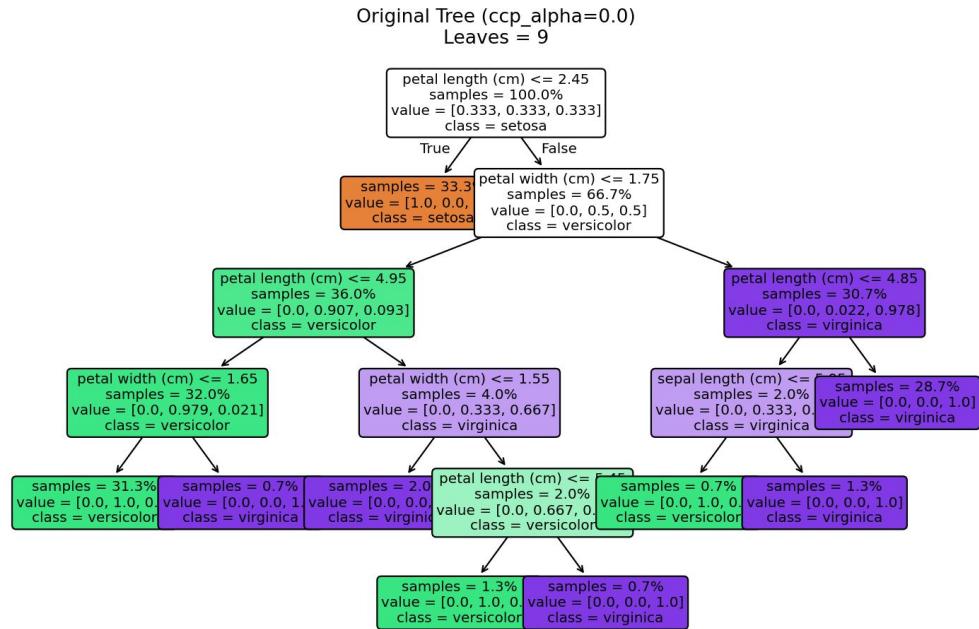




Original Tree



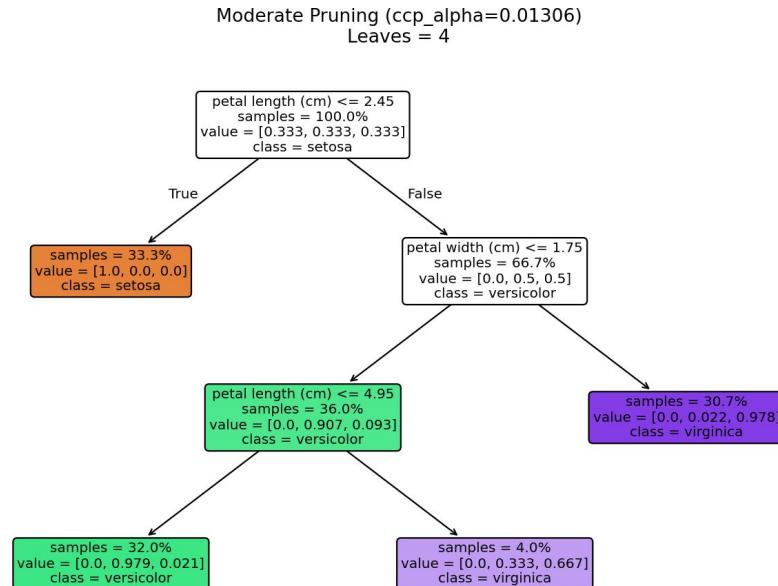
- Large tree grown fully
- Perfect fit to training data
- Risk: overfitting





Moderate Pruning

- Remove weak branches
- Tree becomes smaller
- Slight accuracy loss on train,
better test generalization



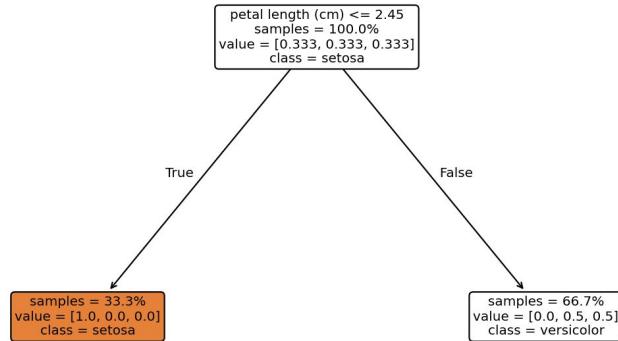


Strong Pruning



Strong Pruning (ccp_alpha=0.2598)
Leaves = 2

- Aggressive reduction of branches
- Tree is much simpler
- Easier to interpret





Decision Trees: Pros & Cons



✓ Pros:

- Easy to understand & explain
- Works with tabular data
- Captures nonlinear interactions

✗ Cons:

- Prone to overfitting (high variance)
- Small changes in data → very different tree
- Weak standalone performance (needs ensembles)





Ensembles



Why Combine Models? 🤝

Single decision tree = unstable, high variance

Small changes in data → very different trees



Idea: aggregate multiple trees to get **more robust predictions**

Analogy:  → 



Bagging = Bootstrap Aggregation



Train multiple models on bootstrap samples (resampling with replacement)

Average predictions (regression) or majority vote (classification)

Reduces variance without increasing bias much



Mathematical View



For regression: average across K models

$$f^{\text{bag}}(x) = \frac{1}{K} \sum_{k=1}^K f_k(x)$$

For classification: majority vote

Works best with unstable models (like decision trees )



Random Forest = Bagging + Random Features



At each split, we **don't look at all features**, only a random few

This makes each tree different → **expert on some small domains**

When we average many such trees → predictions are more stable

That's why Random Forest is a strong baseline in practice





Simple Decision Tree Split Revisit

```
func choose_best_split(data) → (feature, threshold):
    best_split = None; best_gain = -inf
    for feature in features:
        for threshold in possible_thresholds(feature):
            left, right = split(data, feature, threshold)
            if left or right is empty: continue
            gain = impurity(parent) - weighted_impurity(left, right)
            if gain > best_gain:
                best_gain = gain; best_split = (feature, threshold)
    return best_split
```



Random Forest Pros & Cons

✓ Pros:

- Great out-of-the-box performance
- Handles tabular data well
- Robust to overfitting

✗ Cons:

- Slower than single trees
- Less interpretable
- Still limited for very high-dimensional sparse data



Boosting = Sequential Learning



Train models sequentially

Each new model focuses on errors of previous ones

Combines many weak learners → strong learner

Analogy: Coach correcting mistakes **step by step**



Idea of Gradient Boosting



Build trees sequentially, each new tree fixes errors of the previous ones.

Start with a simple model (seed prediction).

Each step: predict **residuals/errors**, fit a new weak learner.

Final model = weighted sum of all trees.



XGBoost / LightGBM ⚡

Engineering optimizations for speed and scalability

Key ideas:

- Regularization (to prevent overfitting)
- Handling missing values
- Parallel training

Popular in Kaggle competitions 🏆



Intuition



Think of a student learning over time:

- Day 1 → learns basics, makes many mistakes.
- Day 2 → focuses on yesterday's mistakes.
- Day 3 → focuses again on remaining mistakes.

Step by step, performance improves.

Gradient Boosting = Residual Learning + Gradient Descent



Key Features & Benefits



- **Learning rate**: controls step size, prevents overfitting.
- **Number of trees**: more trees = lower bias, risk of overfit.
- **Depth of trees**: shallow = weak learners (better generalization).
- **Regularization**: subsampling, shrinkage, pruning.

Boosting vs Bagging:

- **Boosting** = sequential, reduce bias.
- **Bagging** = parallel, reduce variance.



Decision Trees vs Ensembles



Method	Strengths ✓	Weaknesses ✗
Decision Tree 	Simple, interpretable, fast	Overfits, unstable
Bagging 	Reduces variance, robust	Less interpretable
Random Forest 	Strong baseline, feature importance	Slower, less transparent
Boosting 	High accuracy, flexible losses	Sensitive to parameters, less interpretable



Trees vs Deep Learning



Tree excels at tabular data (structured, mixed features)

Deep learning shines at unstructured data (images, text, audio)

In practice → ensembles wins on Kaggle tabular competitions



Rule of thumb



- Tabular → Random Forest / XGBoost
- Image/Text → Neural Networks



Industry Case



The "Unfair" Match-Up



Why does the world's most advanced AI struggle with a simple Excel spreadsheet?

Say, a major NYC fintech company needs to predict loan defaults. The dataset is a classic relational database table containing heterogeneous features:

- Age (Integer), Income (Continuous), Employment Type (Categorical) etc.

The Contenders:

- 🤖 Deep Neural Network (DNN): 100 layers, millions of parameters, trained on massive GPU clusters
- 🎄 XGBoost Ensemble: 500 relatively shallow trees, trained on a standard laptop CPU



Intuition: Smooth Gradients vs. Hard Rules

It is all about the nature of the data and the "**Inductive Bias**"

Neural Networks Expect "Smoothness":

- NNs are essentially massive composite functions optimized by gradient descent
- They thrive on homogeneous data where spatial/temporal proximity implies **mathematical similarity**
- Tabular Data is "Jagged" & Unstructured
 - What is the mathematical "distance" between the categories Teacher and Engineer
 - A Zip Code of 10027 is numerically close to 10028, but demographically completely different

The Struggle: NN waste huge amounts of representational power trying to force these **irregular, completely independent columns into a smooth, continuous geometric space (a manifold)**



The Verdict: Trees are "Tabular Native" 🌲

Tree ensembles **naturally mirror** the logic of structured human data

Rule-Based Decision Boundaries:

- Trees don't care about mathematical distance or gradient flow
- They simply ask orthogonal, hard questions: "Is Income > \$50k?" AND "Is Profession == Teacher?"
- This perfectly matches the conditional logic of tabular data

Built-in Feature Selection:

- Tabular data often contains noisy, irrelevant columns (e.g., user IDs)
- While a Neural Net might overfit to this noise, a Tree model naturally ignores features that don't reduce Information Gini Impurity

The Takeaway:

For tabular data, Feature Engineering + Tree Ensembles > Raw Compute + Deep Learning



Summary



Decision Trees → splitting, pruning

Ensembles → Bagging, Random Forests, Boosting

Takeaway: ensembles make weak learners strong 💪