

COMS4995W32

Applied Machine Learning

Dr. Spencer W. Luo

Columbia University | Fall 2025



From Models to Systems: The Evolution of Applied Machine Learning



Course Journey at a Glance



What we covered (Weeks 1 → 12):

- ML problems → Metrics-driven learning
- Data & features → Linear models
- Decision trees → Ensembles
- Deep learning → Representation learning
- Sequence models → Transformers
- LLMs → RLHF and Agentic systems

Path: From **(un)supervised learning** to **LLM reasoning** and **agentic intelligence**



What This Course Taught Us 🤔

How to formulate real-world problems as learning tasks

- Define objectives, data requirements, and measurable success criteria
- Translate business / scientific requirements into ML formulations

How to build end-to-end ML pipelines

- Data cleaning, feature extraction, model training, evaluation, and deployment
- Automate feedback loops and monitoring in production



What This Course Taught Us 🤔

How to **design experiments and evaluate** appropriately

- Train/validation/test splits, cross-validation, ablation studies
- Understand metrics: precision, recall, F1, AUCPR, perplexity

How to reason about **trade-offs**

- Bias-variance, underfitting vs overfitting
- Scaling vs cost (compute, labeling, iteration speed)



The Core ML Loop



The ML flywheel:

Data → Feature Engineering → Model → Loss → Optimization → Evaluation

- The fundamental feedback loop underlying every ML system
- Each stage influences and constrains the next - improving one improves all



The Core ML Loop



Practical Takeaways

- **Always visualize** intermediate results (train vs validation loss etc.)
- Apply regularization, dropout, or early stopping to balance fit and generalization
- Periodically re-evaluate when data distributions shift (“data drift”)




From Linear to Non-Linear Models

Linear/Logistic regression: simple, interpretable


Decision trees: handle non-linearity, feature interactions

Non-linear models unlock representation power *beyond linear boundaries*

Linear Problem



Nonlinear Problem






Ensembles: The Wisdom of Crowds

Bagging (Random Forest): reduce variance via many weak learners


Boosting (XGBoost): sequentially correct errors, improve performance

Mixture-of-experts: MoE to scale up Transformers


Shallow Tree (max_depth=1)
Train acc = 0.667



Deeper Tree (max_depth=5)
Train acc = 0.993



setosa versicolor virginica






Deep Learning: Representation Learning

Forward pass

Loss computation

Backpropagation




Weight Matrix $W^{(1)}$ (2x4):

$$\begin{bmatrix} 0.5 & -0.3 & 0.8 & 0.4 \\ 0.2 & 0.6 & -0.4 & 0.7 \end{bmatrix}$$

Bias $b^{(1)}$:

$$[0.1, -0.2, 0.3, 0.1]$$



$$\frac{\partial L}{\partial z_i^{(2)}} = \sum_j \frac{\partial L}{\partial p_j} \cdot \frac{\partial p_j}{\partial z_i^{(2)}}$$

$$= \sum_j \left(-\frac{y_j}{p_j} \right) \cdot \frac{\partial p_j}{\partial z_i^{(2)}}$$


Softmax derivative:


if $i = j$: $\frac{\partial p_i}{\partial z_i} = p_i(1 - p_i)$

if $i \neq j$: $\frac{\partial p_j}{\partial z_i} = -p_i p_j$

Final result:

$$\frac{\partial L}{\partial z_i^{(2)}} = p_i - y_i$$







Attention is All You Need

Decoder-only

causal self-attention







Pre-training on Everything

Self-supervised pre-training:

- Train on massive text corpora by predicting masked or next tokens

What it learns:

- Syntax, semantics, reasoning patterns, and broad world knowledge

Outcome:

- General-purpose representations adaptable to many tasks

Impact:

- Pre-training builds the foundation; fine-tuning shapes the purpose



Supervised Fine-Tuning

Converts a pretrained foundation model into a task-specific LLM

Key process:

- Train on curated instruction-reasoning-answer triplets

Data quality > quantity:

- Clean, diverse, well-structured prompts answers yield stronger generalization

Purpose:

- Adapt broad “general intelligence” into practical skills - reasoning for the most



Alignment through RLxF


RLxF = Reinforcement Learning from X Feedback

- Uses **human preference** + **reward models** + **verifiable rewards** to guide behavior

Process:

- 1 Humans/LLMs rank outputs
- 2 Reward model learns verifiable rewards/preferences
- 3 LLM optimized via PPO to maximize those signals

Goal: Align models with verifiable/human preference signals 🏅






PPO: Proximal Policy Optimization ✨

Core idea: small, controlled policy updates

MAXIMIZE objective (simplified version):

$$L(\theta) = \mathbb{E}_t \left[\text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) A_t \right] - \beta \text{KL}(\pi_\theta(a_t | s_t) \| \pi_{\text{ref}}(a_t | s_t))$$

where $r_t(\theta) = \frac{\pi_\theta(a_t | s_t)}{\pi_{\text{old}}(a_t | s_t)}$





What is Next?

Smaller, faster, cheaper LLMs with competitive performance

Agentic application everywhere

Multimodality: integrating text, vision, audio, action...

Advanced alignment: RLHF → RLAIF → Verifiable Rewards → 🤖

New architectures beyond Transformers



Skills to build in the Age of AI

Short term (to landing a job): Strong Python/C++ (vibe debugging/coding), solid ML fundamentals, hands-on project experience, and the ability to clearly **COMMUNICATE** ideas and work with others

Long term (to stay competitive): Develop rapid learning skills, the ability to build and debug quickly with AI, and cross-domain understanding to leverage AI for fast prototyping and real-world impact - then share your lessons to establish your brand and reputation

Mindset: Aim to become a full-stack problem solver - someone who can use AI tools to build, analyze, and deploy ideas **FASTER** 🚀 than before, and learn from success and failures

